

Open-source Technologies and Stream Mining joint Project Documentation

Table of contents

1	Introduction and Background	3
2	Dataset	3
3	System Architecture	4
4	Modeling and Predictions	5
5	Distribution Estimation	8
6	Trend Analysis	8
7	Experiments and Testing	9
8	Appendix	11
9	References	16

Smart City Air Quality Monitoring with Real-Time Stream Analytics (SCAir-IoT)

Abstract

As smart cities deploy large networks of IoT sensors, the ability to collect, process, store, and interpret unbounded environmental data becomes essential for supporting timely decision-making. In this work, a system that processes IoT sensor data deployed in an Italian city was developed with the objectives to detect abnormal changes in the sensor readings in real-time and to create short-term forecasts to support city planning and public health responses. The sensor data related to air quality was simulated, streamlined, processed, modeled, stored, and visualized to develop a full-stack stream mining data science project powered by open-source technologies. The experiments and evaluation of different methods showed that traditional batch-based approaches fail to recognize changes in patterns such as trends, seasonality or in distinguishing actual anomalies from sudden yet expected spikes and drops.

Keywords: *Anomaly detection, Forecasting, Stream mining, Open-source technologies*

Authors

Albazzal, Houmam

Balogh, Máté

Földvári, Ádám

Lahmar, Abderraouf

Nagy, Zsuzsanna

Mentored by

Loubna Seddiki

1 Introduction and Background

In stream mining, we are limited to a portion of data and make decisions real-time in memory. As *Wares, Isaacs, and Elyan (2019)* highlight, in traditional machine learning contexts, data is referred to as batch data which can be loaded into memory in its entirety. According to the authors,

“this is of stark contrast to stream mining, where data streams produce elements in a sequential, continuous fashion, and may also be impermanent, or transient, in nature ... This means stream data may only be available for a short time.”

The authors refer to *Babcock et al. (2002)*, highlighting that

“once an element from a data stream has been processed, it is discarded or archived. It cannot be retrieved easily unless it is explicitly stored in memory, which is small relative to the size of data streams.”

This nature of streaming systems calls for solutions that support the learning from the unbounded data without relying on stored records for every single measurement. Architectural solutions such as *Kafka*, *Spark*, data storages such as *Influxdb*, and algorithms that iteratively learn from the data are important elements in this context.

2 Dataset

The dataset is the *UCI Air Quality* dataset (*Vito, S. (2008)*) which includes responses of gas sensor devices deployed in an Italian city. Besides these device readings, each gas measurement has a counterpart feature which denotes the gas concentration recorded by a co-located certified analyzer. Additionally, readings related to temperature along with absolute and relative humidity are included in the dataset. The records span 1 year from March 2004 to February 2005, and present hourly aggregated measurements. Missing values are denoted with the value of -200.

3 System Architecture




Figure 1: High-level view of the architecture with the utilized open-source technologies denoted for each component.

The stream mining pipeline includes components of *simulator*, *producer*, *consumer*, *data store*, and *dashboards* by which the raw `csv` dataset file was ingested. We utilized open-source technologies of *Java*, *Kafka*, *PySpark*, *Influxdb*, *Streamlit*, and *Grafana* provisioned in a containerized environment via *Docker*. The responsibility of each component is summarized as follows:

Raw data: Static file containing sensor measurements related to air quality.

Sensor data simulator: Reads the raw data file and simulates flow of sensor data.

Kafka producer: Streamlines the simulated sensor data into Kafka topics in their datetime order. *Apache Kafka* serves as a distributed event streaming platform that provides high-throughput, low-latency data feeds with built-in fault tolerance through data replication across brokers (*Apache Kafka Documentation*). The producer component, implemented in *Java*, publishes sensor measurements to topic-partitioned streams, ensuring that each sensor type (e.g., CO, NOx, temperature) is routed to its dedicated topic. This design enables parallel processing and maintains temporal ordering of measurements through the use of datetime-based keys, which guarantees that messages with the same timestamp are processed consistently across the distributed system.

PySpark consumer: Listens to the streamlined data and creates mini batches to call analytical functions—such as Anomaly Detection and Forecasting—on this windowed data. The consumer leverages *PySpark Structured Streaming*, which provides a high-level API for processing continuous data streams with exactly-once semantics and fault tolerance (*Apache Spark Documentation*). The integration between Kafka and PySpark is achieved through Spark’s native Kafka connector, which allows seamless reading from multiple Kafka topics in parallel. The consumer implements a sliding window mechanism that buffers incoming messages per sensor topic (default window size of 8 samples) and triggers analytical computations whenever a new batch arrives. This micro-batch processing model, as described by Zaharia *et al.* (2016), enables real-time analytics while maintaining the benefits of batch processing for complex machine learning operations.

Influxdb: Data is then persisted in the database including the online predictions and the original incoming data.

Streamlit: Streamlit is used to create dashboard prototypes without the necessity of a database storage and connection. Kafka messages are directly consumed by this component to display simple line charts and anomaly detection related information and alerts.

Grafana: Dashboard visualization component that periodically fetches the database for new data to show the latest insights in real-time.

4 Modeling and Predictions

4.1 Overview




Figure 2: High-level view of the offline and online machine learning pipeline.

4.2 Anomaly detection

For *Anomaly Detection*, different techniques were utilized ranging from recognizing sensor readings with missing measurement to detecting local and global outliers, in each sensor separately.

Both, local and global approaches were developed with the general novelty, outlier detection methodology (*Müller and Chiu (2024)*) in mind, where the steps are:

1. Transform the data, to better reflect properties of interest
2. Obtain a novelty function
3. Apply peak picking algorithm on the novelty function

To detect local patterns in form of sudden changes (peaks and valleys), an 8-sample window was utilized on the streamlined data. As part of the data transformation, values of -200 were replaced with NA. The most recent element in the window was selected as a test sample to make predictions for, while the rest of the in-window samples (at most 7) were designated as historical training points. Derivative operation was applied on the window data to turn a measurement value x_T into the deviation from x_{T-1} . This resulted in the novelty function, where NA values were first imputed with the window median but were transformed back to NA once the difference values were obtained. During the “training”, the in-window estimator calculates statistics (IQR , $Q1$, $Q3$) from the in-window historical samples to perform a traditional outlier detection test using the $Q1 - 1.5 \times IQR$ and $Q3 + 1.5 \times IQR$ as thresholds to flag the test sample as outlier.

To account for out-of-window global patterns and to implement a detector that is applicable in stream mining systems where we cannot always rely on storing every incoming data point, an online detector was implemented using the *T-digest* algorithm. This method uses the same peak picking strategy using the traditional outlier thresholds, yet the quartiles along with the *Interquartile Range* get iteratively updated by each new sample. The algorithm maintains a centroid-based representation about these statistics using the incoming data and updates its state through merges.

4.3 Offline Forecasting

The offline forecasting module represents a traditional Supervised Machine Learning pipeline where models are trained on historical datasets and deployed for inference on streaming data. This approach prioritizes model complexity and accuracy over immediate adaptability, assuming that the underlying statistical properties of the air quality data remain relatively stable over short periods.

4.3.1 Model Training and Architecture

The foundation of this module is established in `test_Modelling.ipynb`. The training process begins by loading the UCI Air Quality dataset and performing rigorous preprocessing, including datetime indexing and handling missing values (mapping sensor error codes like **-200** to **NaN**). The core of the strategy relies on **supervised feature engineering**, where the time-series problem is transformed into a regression problem. The system constructs a rich feature set comprising **lagged values** (1, 2, 3, 6, 12, 24 hours) to capture immediate and daily dependencies, **rolling statistics** (mean and standard deviation over 3, 6, 12-hour windows) to smooth volatility, and **temporal embeddings** (hour of day, day of week, month).

The primary algorithm utilized is the **Histogram-based Gradient Boosting Regressor** (HGBR). This algorithm is chosen for its efficiency with large datasets and native handling of missing values (NaNs), which are common in sensor networks. Separate models are trained for specific forecast horizons (H+1, H+2, and H+3 hours ahead). These trained models, along with their metadata, are serialized using `joblib` and stored in the `artifacts/` directory, ensuring that the inference engine has access to pre-validated statistical patterns.

4.3.2 Inference Pipeline and State Management

The inference logic is encapsulated in `offline_forcasting/offline_forecasting.py` within the **OfflineForecaster** class. This class acts as a singleton **Sliding Window Listener** that consumes raw sensor data arriving from the Simulator via Kafka. A critical challenge in offline-to-online deployment is **feature consistency**; the inference engine must reconstruct the exact feature set used during training. To achieve this, the forecaster maintains in-memory buffers (`deque`s) for every sensor topic.

As data streams in, the **OfflineForecaster** aligns the history across different sensors. Once the buffer accumulates enough samples to calculate the required lags and rolling windows (left-padding with **NaNs** during cold starts), it constructs a single-row feature vector. This vector is passed to the loaded HGBR pipelines. To ensure robustness, the system respects the `feature_names_in_` attribute of the saved models, preventing column mismatch errors. Predictions are generated both **on-demand** (triggered by window updates) and via a **periodic background thread** (defaulting to a 10-second cadence). Finally, the forecasted values are persisted to **InfluxDB** via the `dbWriter`, allowing for immediate visualization and comparison against actual incoming values.

Suggestions for Future Improvements

- **Automated Retraining Pipeline:** Currently, models are static. Implementing a “Champion/Challenger” system that periodically retrains models on the most recent week of data and swaps them in production would mitigate long-term model rot.
- **Deep Learning Architectures:** Replacing Gradient Boosting with LSTM (Long Short-Term Memory) or Transformer-based architectures (like Temporal Fusion Transformers) could better capture long-range dependencies, provided the inference latency remains acceptable.
- **Feature Store Integration:** Decoupling feature calculation from the application logic using a Feature Store (e.g., Feast) would ensure stronger consistency between the training notebook and the inference script.

4.4 Online Forecasting

The **Online Forecasting** module, implemented in `online_forecasting/online_forecasting.py`, addresses the limitations of static models by implementing **Incremental Learning**. Unlike the offline approach, this system does not rely on pre-existing artifacts. Instead, it initializes fresh models that learn continuously from the data stream, allowing the system to adapt rapidly to **concept drift** (e.g., sudden changes in sensor calibration or environmental conditions).

4.4.1 Theoretical Context and Algorithm

Online learning differs fundamentally from batch learning by performing a “predict-then-update” cycle. The system utilizes the **Stochastic Gradient Descent (SGD) Regressor**, a linear model optimized for streaming. While linear models are generally less complex than tree-based ensembles, SGD is computationally inexpensive and capable of updating weights sample-by-sample (`partial_fit`).

The workflow is strictly sequential:

1. **Extract Features** (X_t): Current lags and covariates are computed.
2. **Predict** (y_{t+1}): The current model predicts the next step.
3. **Wait**: The system waits for the actual observation of y_{t+1} .
4. **Update**: Once the actual value arrives, the model calculates the error (loss) and updates its weights via backpropagation to minimize future errors.

Crucially, because SGD is sensitive to feature scaling, the pipeline includes an incremental **StandardScaler** that also updates its mean and variance estimates on the fly, ensuring that the gradient descent converges correctly even as the statistical properties of the raw data shift.

4.4.2 Adaptive Pipeline and Robustness

The **OnlineForecaster** manages its own bounded buffers (defaulting to 200 points) to construct features similar to the offline model but optimized for speed: short-term lags (1, 2, 3, 6 hours) and covariates (Temperature and Humidity interactions). To prevent the model from making erratic predictions during the initial “warm-up” phase or during sensor malfunctions, several **robustness mechanisms** are implemented.

The system maintains streaming **Quantile Sketches** (using T-Digest concepts) and calculates **Z-scores** on a short sliding window. If an incoming value is detected as a statistical anomaly (e.g., $> 3.5\sigma$), the system flags it. Furthermore, predictions are **clamped** to a dynamic confidence band derived from recent rolling statistics. If the model produces a prediction that is physically implausible or diverges wildly from the recent trend, the system falls back to a simple rolling mean. This hybrid approach ensures that the forecaster remains stable even when the SGD model is under-fitted or facing outlier data. The resulting one-step-ahead forecasts are written to **InfluxDB**, tagged specifically as `online_pred`.

4.4.3 Suggestions for Future Improvements

- **Adaptive Learning Rates**: Implementing mechanisms to dynamically adjust the SGD learning rate (`eta0`) based on the volatility of the error signal. If drift is high, the learning rate should increase to adapt faster; if stable, it should decrease to converge.
- **Non-Linear Online Models**: Moving beyond linear SGD to non-linear online algorithms, such as **Hoeffding Adaptive Trees** or Kernel-based Recursive Least Squares (KRLS), to capture complex relationships without the full overhead of batch retraining.
- **Ensembling**: Creating a weighted ensemble of several online models with different learning rates (e.g., one “fast” learner and one “slow” learner) to balance stability and adaptability.

5 Distribution Estimation

To efficiently estimate the distribution of streaming data, we use **t-digest** *Dunning, T., Ertl, O. (2021)*, an online algorithm optimized for accurate quantile estimation with limited memory. Instead of storing all incoming observations, t-digest compresses the data into a set of **centroids**, each representing a cluster of nearby values.

t-digest is governed by a **compression parameter** δ , which controls the maximum number of centroids maintained in the structure. A higher compression value results in more centroids and therefore higher quantile accuracy—particularly near the tails of the distribution—at the cost of increased memory usage. Lower compression values produce a more compact summary but with reduced precision.

As new data points arrive, t-digest incrementally updates or merges centroids according to the compression constraint, ensuring the digest remains small and efficient. This allows real-time estimation of quantiles and other distributional properties without storing the entire dataset.

t-digest is thus well-suited for monitoring evolving data streams, detecting changes in distribution behavior, and enabling data-driven decisions in systems where memory and computation are constrained.

6 Trend Analysis

For trend analysis, we adopt an approximative variant of the **Mann–Kendall (MK) test** *Kamal, N., & Pachauri, S. (n.d.)*, a **non-parametric statistical test** used to detect the presence of a **monotonic trend** (increasing or decreasing) in a time series.

6.1 Hypotheses

Hypothesis	Description
H_0 (Null)	There is no monotonic trend in the time series.
H_a (Alternative)	One of the following: 1. Upward trend 2. Downward trend 3. Either trend

6.2 Why Mann–Kendall?

- **Robust to missing data:** Missing points only reduce the sample size, which may affect significance but does not invalidate the test. In our data, some features have many missing measurements.
- **Distribution free:** The test does not assume any specific probability distribution, which is unknown in streaming data.

6.3 Online Mann–Kendall Adaptation

6.3.1 Approximating the Sign Function

Traditional MK computes:

$$S = \sum_{i < j} \text{sgn}(x_j - x_i)$$

In our online version:

- The sign function is approximated using the CDF of the current data, maintained via a t-digest.
- Quantile information yields an approximate ordering. For example, a value at the 7% quantile is likely smaller than one at the 15% quantile.

Approximation:

$$\text{sgn}(x_j - x_i) \approx 2 \cdot \text{CDF}(x_j) - 1$$

This reduces computation and memory while preserving the ordering structure.

6.3.2 Handling Ties

Sensor data resides in continuous space, and features vary in scale. To adapt MK for this scenario:

- A dynamic tolerance ϵ_{tie} based on the moving range ensures robust tie handling:

$$\epsilon_{\text{tie}} = \max(\text{moving range} \cdot \text{rel_tol}, 10^{-8})$$

- Tie counts adjust the variance term. A lossy counting *Manku, G. S., & Motwani, R. (2002)* strategy tracks frequencies efficiently.

6.3.3 Variance and Trend Significance

Variance of S accounting for ties:

$$Var(S) = \frac{n(n-1)(2n+5) - \sum t(t-1)(2t+5)}{18}$$

Real-time Z-score:

$$Z = \begin{cases} (S-1)/\sqrt{Var(S)}, & \text{if } S > 0 \\ (S+1)/\sqrt{Var(S)}, & \text{if } S < 0 \\ 0, & \text{if } S = 0 \end{cases}$$

Trend classification:

- Significant increasing trend
- Significant decreasing trend
- No significant trend

7 Experiments and Testing

High-level and sensor-wise dashboards were created in *Grafana*, while early prototypes were created in *Streamlit* (Figure 3).

7.1 Anomaly detection

For anomaly detection, markers with different colors were used for the global and local abnormal detections, along with the highlights of missing values, plotted over the actual sensor readings. Examples of this dashboard is displayed in Figure 4 - Figure 5.

The dataset did not contain ground truth for anomalies. To see how these detectors perform in terms of positive predictive capability, the following manual labeling process was followed:

1. Designate the rush hour segment (6-10 AM) in each day in an entire month from the static dataset
2. Combine the sensor readings in those segments
3. Apply outlier detection (LOF) on this subset of data to obtain observations that occur rarely and does not align with normal patterns of the segments
4. Label these observations as anomalies
5. Evaluate whether local and global abnormality detectors can recall these special anomalies

As we can see in Figure 6, the in-window and global abnormality detectors were able to recall only a few of these special anomalies and predict substantially more cases as abnormal. One reason might be that these detectors cannot recognize the seasonality, the reoccurring peaks within a segment of each day throughout a month. A refinement may involve the utilization of more advanced anomaly detection algorithms or creating separate global models for each segment of interest.

7.2 Forecasting

The output of online forecasting was plotted along with a live performance evaluation panel including metrics of *Absolute Error (AE)*, *Cumulative Mean Absolute Error (MAE)*, *Squared Error (SE)*, *Cumulative RMSE*, and *Cumulative Bias (mean signed error)* to compare the predicted line against the actual sensor readings in real time. This is displayed in Figure 7, while comparison of offline and online forecasting is shown in Figure 8.

7.3 Trend Analysis

We applied the Mann–Kendall trend detector to visualize the types of trends present in the data stream. These trend indicators were plotted together with key distributional statistics, such as quantiles, mean, and variance, to provide a comprehensive overview of the data behavior. When a trend occurs, distributional characteristics often evolve: the mean may shift, the variance may increase or decrease depending on the trend’s direction and magnitude, and the quantiles may move, indicating systematic changes across different portions of the distribution. The resulting dashboard supports real-time tracking of these distributional shifts and facilitates the early detection of potential trends within the data stream. The related dashboard is depicted in Figure 9.

7.4 Overview Dashboards

This dashboard acts as the primary observability layer for the SCAir-IoT project, aggregating data from the entire sensor network to provide high-level situational awareness.

For real-time system status, the System Overview panel utilizes gauge visualizations to display the most recent readings. It includes environmental factors and key pollutants, allowing users to instantly assess if current values are within expected physical ranges.

The Anomaly Detection & Monitoring section visualizes the outputs of the stream mining algorithms through two key components. A summary panel, ‘Sensors by Anomaly Count’, highlights the total number of anomalies detected per sensor in the given time window. Alongside this, the ‘Anomaly Matrix’ serves as a state timeline that temporally tracks the two distinct detection strategies described in the modelling section: “global” outliers and “local” in-window spikes. Red blocks within this heatmap indicate detected anomalies.

The Global Trends section features a multi-line time series chart that aggregates the mean values of all sensors, visualizing daily pollution cycles and long-term trends.

The correlation dashboards are helping to analyze physical and chemical relationships to validate sensor consistency: The “Weather Dynamics” chart reveals an inverse pattern where relative humidity drops as temperature rises. This confirms the meteorological principle that warmer air

holds more water vapor. Significant divergences from this trend can serve as early indicators for potential sensor faults.

The scatter plot demonstrates a strong linear relationship between Carbon Monoxide (CO) and Nitrogen Oxides (NOx). Their values are rising and falling simultaneously. This clustering confirms a shared emission source, as both pollutants are primary byproducts of combustion engines and fluctuate instantly with traffic flow.

The Single Sensor Deep Dive dashboard provides an analysis of individual data streams, allowing users to dynamically select specific sensors via the Topic variable. It focuses on validating the statistical modelling and anomaly detection logic.

The dashboard also visualizes the stochastic properties of the data stream. It compares actual readings against calculated Moving Averages and Standard Deviation bands (upper/lower bounds) to visually identify outliers. A Data Distribution chart plots the current value against dynamic quantiles (Min, Q1, Median, Q3, Max), offering a visual representation of the T-digest and IQR-based algorithms used for outlier detection.

The Sensor Health Status timeline tracks specific error states, differentiating between “Data Missing” events and algorithmic “Anomaly Status” flags, while a Frequency Histogram analyzes the distribution of incoming values to detect skewness or sensor drift.

The related dashboards are depicted in Figure 10, Figure 11, and Figure 12.

8 Appendix




Figure 3: Anomaly detector prototype dashboard implemented in *Streamlit*.




Figure 4: Abnormality detection dashboard.



Figure 5: Abnormality detection dashboard (missing values).



Figure 6: Abnormality detection dashboard.



Figure 7: Predictions and performance evaluation dashboard for forecasting.



Figure 8: Predictions and performance evaluation dashboard for forecasting, based on both offline and online models.



Figure 9: Trend analysis dashboard.



Figure 10: Overview Dashboard for all sensors I.




Figure 11: Overview Dashboard for all sensors II.



Figure 12: Overview Dashboard for a single sensor

9 References

- Wares, S., Isaacs, J. and Elyan, E. (2019). Data stream mining: methods and challenges for handling concept drift. *SN Applied Sciences*, 1(11), 1412.
- Babcock, B., Babu, S., Datar, M., Motwani, R. and Widom, J. (2002, June). Models and issues in data stream systems. In Proceedings of the twenty-first ACM SIGMOD-SIGACT-SIGART symposium on Principles of database systems (pp. 1–16).
- Vito, S. (2008) Air Quality Dataset. UCI Machine Learning Repository. Available at: <https://archive.ics.uci.edu/ml/datasets/Air+Quality> (Accessed: 25 November 2025).
- Müller, M. and Chiu, C.Y., 2024. A basic tutorial on novelty and activation functions for music signal processing. *Transactions of the International Society for Music Information Retrieval*, 7(1).
- Apache Kafka Documentation. Available at: <https://kafka.apache.org/documentation/> (Accessed: 25 November 2025).
- Apache Spark Documentation. Structured Streaming Programming Guide. Available at: <https://spark.apache.org/docs/latest/streaming/index.html> (Accessed: 25 November 2025).
- Zaharia, M., Das, T., Li, H., Hunter, T., Shenker, S., & Stoica, I. (2016). Discretized streams: Fault-tolerant streaming computation at scale. In Proceedings of the 24th ACM Symposium on Operating Systems Principles (pp. 423-438).
- Kamal, N., & Pachauri, S. (n.d.). Mann-Kendall Test – A Novel Approach for Statistical Trend Analysis. Faculty of Computing and Information Technology, Himalayan University, Arunachal Pradesh; Department of CSE/IT, IIIMT College of Engineering, Greater Noida, U.P.
- Dunning, T., & Ertl, O. (2021). Computing extremely accurate quantiles using t-digests. MapR Technologies, Inc., Santa Clara, CA; Dynatrace, Linz, Austria.
- Manku, G. S., & Motwani, R. (2002). Approximate frequency counts over data streams. Stanford University.